

Image Formation Model:
$$I(x,y) = i(x,y) \cdot r(x,y)$$

 $i(x,y)$ - Amount of source illumination incident on the scene, $0 \leq i(x,y) < \infty$
 $r(x,y)$ - Amount of illumination reflected by objects in the scene, $0 \leq r(x,y) \leq 1$

- 0 for total absorption, 1 for total reflectance (All energy is reflected)

Resolution:
1. Spatial Resolution
Higher Resolution \rightarrow Will have a detailed look of what we are looking at
Lower Resolution \rightarrow Very pixelated and we will not know what we are looking at
2. Intensity Resolution
More Bits \rightarrow More grey levels and can discern between different levels
File Size: Spatial and Intensity resolution determines the file size
$$B = M \cdot N \cdot K$$

 $M \cdot N$ - Size of the image, K - Bit depth of the image
Maximum number of gray scale levels: $(N - 1) \times n + 1$, Where N is the number of grey levels and n is the number of channels that it is split equally by

Filtering
Image Normalization (Whitening): Resulting image pixels are zero-mean, unit variance.

- Allows for removal of contrast and constant additive luminance variations

$$\mu = \frac{\sum_{i=1}^N \sum_{j=1}^M p_{ij}}{IJ}, \quad \sigma^2 = \frac{\sum_{i=1}^N \sum_{j=1}^M (p_{ij} - \mu)^2}{IJ}, \quad x_{ij} = \frac{p_{ij} - \mu}{\sigma}$$
Gamma Mapping: Mappings can be non-linear
$$x_{ij} = 255 \cdot \left(\frac{p_{ij}}{255} \right)^\gamma, \quad \gamma > 0$$

- Small $\gamma \rightarrow$ Increases the range for the dark areas, makes the dark areas become brighter.
- Large $\gamma \rightarrow$ Increases the range for the bright areas, makes the bright areas become darker.

Histogram Stretching: Stretch the boundaries to $[0, K]$
$$x_{ij} = \frac{(p_{ij} - f_{\min})}{(f_{\max} - f_{\min})} \times K$$

 f_{\min} - Minimum value observed, f_{\max} - Maximum value observed, p_{ij} - Value at pixel ij , K - maximum value that is possible in that image

Histogram Equalisation: Forces the cumulative distribution of the intensity to be linear
$$x_{ij} = K \cdot c_{p_{ij}}$$

 K - Maximum value observed, $c_{p_{ij}}$ - Cumulative distribution value at pixel ij
Note: If the images have a lot of boundaries vs another that does not have much have the same histogram originally. If we do blurring, the one with more boundaries will have more counts for in between values and therefore their histograms will no longer be the same.

Image Normalisation (Whitening) vs. Histogram Equalisation:

- Equalisation boosts contrast more
- However, equalisation is much more expensive

Otsu's Method: Minimise the sum of weighted variance of object and background by choosing the optimal threshold.
$$T^* = \arg \min T_1 \sigma_1^2(T) + w_2(T) \sigma_2^2(T)$$

 $\sigma_1^2(T), \sigma_2^2(T)$ - Variance of pixels less than or equal to and greater than threshold $w_1(T), w_2(T)$ - Number of pixels less than or equal to and greater than threshold

Normalised Cross-Correlation: Normalize cross correlation to reduce effect of white patch
$$x_{ij} = \frac{1}{|F| |w_{ij}|} \sum_{u=-k}^k \sum_{v=-k}^k f_{uv} \cdot p_{i+u, j+v}$$

 $|F|$ - Magnitude of the filter, $|w_{ij}|$ - Magnitude of the input widow size. The remaining part is the normal cross correlation term. The front part is a scaling factor.

Noise in Image:

Noise	Description	Solution
Impulse & salt and pepper noise	Random occurrences of white (and black) pixels, usually due to defective sensor elements	Median Filter
Gaussian Noise	Variations in intensity drawn from a Gaussian distribution due to inherent noise of the sensor	Gaussian Filter, Box Filter

Filters:

Filter	Matrix (Example)	Effect of the Filter
Horizontal Sobel Filter	$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$	Detects horizontal gradients and therefore can detect vertical lines. Takes $O(w \times h)$
Vertical Sobel Filter	$\begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$	Detects vertical gradients and can detect horizontal lines. Takes $O(w \times h)$
Median Filter	3×3 Median Filter. Just take the median value of the 3×3 area	Gives a boxy look. Helps to remove noise from the image. Note that the if the kernel is small enough, we won't have a boxy effect. Removes spikes in the image.
Box Filter	$\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	It is technically just a Gaussian Blur. Need to look at the averaging term to see how much is being averaged
Gaussian Filter	$\frac{1}{16} \begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix}$	Makes use of the neighbouring values and therefore gives a smoothing effect. Could perform a better job than box filter.

Shift Filter	$\begin{bmatrix} 15 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix}$	Shifts the image. Note that it will move in the opposite direction of the filter. (I.e. This filter is up and left, this will cause the image to move down and right)
Sharpen Filter	$-\frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & -17 & 1 \\ 1 & 1 & 1 \end{bmatrix}$	Accentuates the details on the image. Accentuates difference with local average. Noise and edges will get emphasised but no effect on flat regions.
1D Derivative Filter	$\begin{bmatrix} -1 & 0 & 1 \end{bmatrix}$	Detects edges, in this case will be horizontal lines
1D Laplace Filter	$\begin{bmatrix} 1 & -2 & 1 \end{bmatrix}$	Detects edges, in this case will be horizontal lines. Similar to derivative filter
2D Laplace Filter	$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$	Like sharpening filter but helps us to detect edges through accentuating the differences.

Template Matching:

Type	What would be detected	Issues
Cross Correlation	Bright patches will give the strongest response	Due to bright spots having large intensity values , it will give large response for the template match even if it doesn't match
Convolution		Not really an issue: If the patches that we try to detect is not the exact template. Then the one that looks most similar ("blob" that looks the most similar) will give the strongest response.
Normalised Cross Correlation	Allows for better detection with respect to the template. Not as susceptible to bright patches.	

Possible Pre-Processing Steps:

Pre-Processing Step	Benefits	Cons
Sharpening	Can sharpen some of the details and helps with processing after	Too much could cause over emphasis on the noise
Intensity Thresholding - Use Otsu's Method for Automatic Thresholding	Useful for removing extreme background (Black or White)	If our image has bright spots / dark spots as the features then may not be optimal If there are already dark or bright spots (0 / 255 type) then it won't be able to stretch the range
Histogram Stretching	Could help to improve contrast	
Histogram Equalisation	Helps to map the distribution of the histogram to be linear. Helps with improving contrast, probably better than histogram stretching in most cases	May cause some pixels to darken / lighten due to higher proportion of brighter / darker patches overall in the image.
Canny Edge Detection	Helps to convert thick regions to thinner regions for better detection of lines afterwards	

Gradients
Gradients and Edges:

Filter	Description	Benefits / Cons
Horizontal Sobel Filter	Type of derivative filter to detect vertical lines.	Susceptible to noise. However, detected edges will be very sharp.
Vertical Sobel Filter	Type of derivative filter to detect horizontal lines.	Susceptible to noise. However, detected edges will be very sharp.
Derivative of Gaussian (DoG)	Combines the Gaussian Filter with a derivative filter (Like Sobel)	Applies blurring to the image to help remove some of the noise. However, the detected image may not be as sharp. Edges are at peaks of white points.
Laplacian of Gaussian (LoG)	Combines the Gaussian Filter with a Laplacian filter	Applies blurring to the image to help remove some of the noise. However, the detected image may not be as sharp. Blob detector. Edges are at zero-crossing but not a preferred choice since it is hard to find. Can be estimated by Difference of Gaussians

Canny Edge Detector:

- Filter image with Derivative of Gaussian (DoG)
- Find magnitude and orientation of gradient
- Non-maximum suppression
Thin wide "ridges" down to single pixel width
- Linking and thresholding (hysteresis):
Use the high threshold to start edge curves and connect them with the low threshold edges

Lines and Houghs
Double Intercept Form: $x/a + y/b = 1$
Normal Form: $x \cos \theta + y \sin \theta = \rho$, $a = \cos \theta / \rho$, $b = \sin \theta / \rho$
Line (Normal Form) Detection with Hough Transform:

- Quantize Parameter Space (θ, ρ)
- Create Accumulator Array $A(\theta, \rho)$
- Set $A(\theta, \rho) = 0, \forall \theta, \rho$
- For each image edge point (x_i, y_i)
For each element θ
Solve for $\rho = x_i \cos \theta + y_i \sin \theta$
Increment $A(\theta, \rho) = A(\theta, \rho) + 1$

Hough Transform for Circle:
$$(x-a)^2 + (y-b)^2 = r^2$$

Note that we will take r to be a constant and the centre of the circles are the parameters.

We will create a different accumulator for each radius.

Leveraging Gradient Information:

- If we have a lot of points, it could cause a lot of votes to be distributed across the accumulator and making it hard to discern the actual maxima.
- We can make use of the gradient information to **minimise the votes and reduce computation time**.
- Make use of only the 2 directions of the orientation to cast the votes rather than all directions.

Generalized Hough Transform:
Assumption: Translation is the only transformation. (I.e. orientation and scale of the arbitrary shape are fixed)

- Define a model shape by its boundary points p_i and reference point a that we will try to vote towards.
- For each edge point:
 - Use its gradient orientation ϕ to index into stored table.
 - Compute some averaging statistics for each of the points (e.g. mean, mode)
 - Use retrieved r vectors to vote for reference point.

Hough Voting for Object Detection:
Idea: Instead of indexing displacements by gradient orientation, index by matched local patterns (visual codewords)
Note: This visual codewords will be pre-generated and it should be able to discern between different directions. (I.e. For car tyres, we can point to 2 sides and when 2 tyres vote for the same centre, it will be centre of a car)

Hough Transform:

Pros	Cons
All points are processed independently so can cope with occlusions and gaps	Search time complexity increases exponentially with the # model parameters
Some robustness to noise: Noise points unlikely to contribute consistently	Non-target shapes can produce spurious peaks in parameter space
Can detect multiple instances of a model in a single pass For instance, if there are multiple lines in the same image, they can be detected at once	Quantization: Can be tricky to pick a good grid size

Segmentation
K-Means Clustering: Minimises average squared distance within-cluster. It does not try to find the actual cluster means but works on minimising the objective function instead

- Given K , randomly initialize the cluster centers, c_1, \dots, c_k
- Given clusters centers, determine points in each cluster
For each point p_i , find the closest c_j . Put p_i into cluster j
- Given points in each cluster, solve for c_j
Set c_j to be the mean of points in cluster j
- If c_j have changed (up to some threshold), repeat Steps 2, 3.

Pros	Cons
Simple, fast to compute	Clusters based on intensity don't have to be spatially coherent
Converges to local minimum of within-cluster squared error	Requires sufficient number of points to ensure convergence to the correct clusters
	Sensitive to initial centers
	Sensitive to outliers (since we are taking mean)
	Detects only spherical clusters
	Assumes that means can be computed (efficient, meaningful)
	Hard to determine the number of k we need initially

SLIC Superpixeling: A group of pixels that share common characteristics locally
Computation:
 n_{sp} - Number of superpixels, parameter that we set
 n_{tp} - Total number of pixels, determined by $W \times H$
 $s = \left(\frac{n_{tp}}{n_{sp}} \right)^{1/2}$ - Initial spacing of each superpixel.
 $n_W = W/s$ - Number of pixels in each superpixel's width
 $n_H = H/s$ - Number of pixels each superpixel's height
 $D = \left(\left(\frac{d_{xc}}{d_{xm}} \right)^2 + \left(\frac{d_{ys}}{d_{ym}} \right)^2 \right)^{1/2} = \left[d_x^2 + d_y^2 \right]^{1/2}$ - Compose distance and note that s and c will determine how much weightage is placed on each of the terms.
Changing the value of c :
If c large \rightarrow Spatial information more important. We will have more regularly shaped segments. It will be a hexagonal shape which has the shortest distance to the center. Shifts in the centers can be computed through the perpendicular bisector.
If c small \rightarrow Colour information more important. We will have more irregularly shaped segments.
 c - Maximum colour value. s - Maximum spatial value is the spacing between 2 superpixels.

Algorithm: Note that we can do segmentation with this after the super-pixeling \rightarrow Make use of K-Means as well.

- Compute the initial super pixel clusters,
 $m_i = [r_i, g_i, b_i, x_i, y_i]^T, \quad i = 1, 2, \dots, n_{sp}$
Sample the image at regular grid steps, s . Move the cluster centers to the lowest gradient position in a 3×3 neighbourhood. For each pixel location, p , in the image, set a label $L(p) = -1$ and a distance $d(p) = \infty$
 - Note that we do not want to initialise at an edge or noise so that it will be similar to the points around it

- Assign samples to cluster centers. For each cluster center $m_i, i = 1, 2, \dots, n_{sp}$, compute the distance $D_i(p)$ between m_i and each pixel p in a $2s \times 2s$ neighbourhood around m_i . Then for each p and $i = 1, 2, \dots, n_{sp}$, if $D_i < d(p)$ let $d(p) = D_i$ and $L(p) = i$. (Only compute distance to closest set of cluster centers)
- Update the cluster centers by computing the new mean (as per K means)
- Test for convergence:** Compute Euclidean norms of the differences between the mean vectors in the current and previous steps. Check the residual errors for the super pixels, if the sum is less than T , a specific threshold then go to Step 5 else repeat Step 2.
- Post processing (Optional to create "Stained-Glass" Effect): Replace all the super pixels in each region, G_i by their average value m_i

Pros	Cons
Could be faster than K-Means since we are only considering neighbours in a range of $2s \times 2s$	No provision is made to enforce connectivity of super pixels. Could have isolated pixels to remain after convergence
Can be used in other color spaces and distance can be other values that we can derive meaningful distance measure	If we do segmentation, we may have not get the exact shape back. (I.e. circle)

Mean Shift: Find modes or local density maxima in feature space
Algorithm: Do the following for each data point
For each point x_j , where $j = 1, \dots, n$

- Initialise the density window $x: x = x_j$
- Compute the mean shift vector m
$$m_{h,g}(x) = \frac{\sum_{i=1}^n x_i g\left(\left\|\frac{x-x_i}{h}\right\|\right)}{\sum_{i=1}^n g\left(\left\|\frac{x-x_i}{h}\right\|\right)} - x$$
- Shift the density window and update
$$x' = x + m(x), \quad x = x'$$
- Repeat until the centroid stops moving (convergence)

Segmentation with Mean Shift:

- Find features (colour, gradients, textures, etc) to represent each pixel
- Initialize density windows at individual feature points
- Perform mean shift for each point until convergence
- Merge pixels whose feature points that end up near the same mode or "peak" into the same segment

Increasing bandwidth, h / window size: When bandwidth increases, we will be including more points in the computation and makes the mean shift vector more accurate and will try to converge towards the actual mean. **Number of segments generated decreases.**
Note: We need a large enough amount of neighbours to ensure that the algo can progress.

Pros	Cons
Finds variable number of modes (don't have to pre-specify)	Output depends on bandwidth, h and selecting it needs trial and error
No prior assumptions on cluster shape	Computationally expensive and slow to run
1 parameter only (window size / bandwidth) which represents the scale of local max we are intending to find	Scales poorly with feature space dimension since we need to repeat with all points
General algorithm for mode-finding	
Robust to outliers	

Speedups:

- Assume that all points within radius r of end point after mean shift will belong to the same cluster / mode. We can assign them to the same cluster and don't have to run on mean shift on them
- Assume that points in the same path will be attracted to the same basin of attraction. Assign all points within radius c of search path to mode. But it will most likely be smaller than r .

Textures
Choice of window scale for texture representation: Perform scale selection by looking for window scale where statistic (texture representation) does not change much.
Gabor Filters: Can be used to create many different filters for the use of a filter bank
$$f_{mn} = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{m^2 + \gamma n^2}{2\sigma^2}\right) \sin\left[\frac{2\pi(\cos(\omega)m + \sin(\omega)n)}{\lambda} + \phi\right]$$

 m, n - 2D coordinates of the kernel, σ - Controls the spread of the Gaussian kernel, ω - Controls the orientation of the sinusoid in 2D, will influence the orientation of the detected feature, λ - Wavelength of the sinusoid (Tells us the frequency of the sinusoid for us to see the rest of the peaks when merged with the Gaussian), ϕ - Controls the amount of shift from center of the kernel.

Gabor Filter Estimations: If scale is small compared to the frequency, Gabor Filters can be used as approximate derivative operators.

- Gaussian Derivative: Can be estimated by odd Gabor Filter
- LoG: Can be estimated by even Gabor Filter

Computing Textons: Replace each pixel with integer representing texture "type"

- Apply filter bank to (training) images
- Cluster in feature space and store cluster centers. This forms a texton dictionary. Note that this is done using the training images.
- For new (test) image, filter image with same filter bank to get feature representations for each pixel based on feature vectors at each pixel, assign to the nearest cluster. The cluster ID is the texton ID.

Histogram of Textons:
For a given region, compute a histogram of textons as representation: vector storing number of occurrences of each texton.
Note: When we have relative counts, it could help if the scale of the images are different

Make use of Texton Histogram for Classification: If a new region is most similar to a texton histogram → Assign to the label
Drawbacks: Effectiveness is dependent on the size of the window as the window size could make it hard to localise.

Texture Boundaries:

- Consider a disc, split into two halves by a diameter of a particular orientation
- Measure the difference in texture between the two halves by comparing texton histograms
- Try all possible orientations → If we get a strong difference then it should be a texture boundary

Keypoints
Equivariance: If there is a transformation on the image, there should be the same transformation on the keypoint. (Location of the keypoint should be transformed in the same manner as the image)
Invariance: Image is transformed and detection (score) does not change.

Harris Corners: We have a fast approximation for the estimation of the values
 $R = \lambda_1 \lambda_2 - \kappa (\lambda_1 + \lambda_2)^2$, $\det(H) = \lambda_1 \lambda_2$, $\text{trace}(H) = \lambda_1 + \lambda_2$
 λ_{\max} – Eigenvector with direction of larger increase in error → Minor Axis ($\lambda_{\min}^{1/2}$)
 λ_{\min} – Eigenvector with direction of smaller increase in error → Major Axis ($\lambda_{\max}^{1/2}$)
Determinant of 2×2 : $ad - bc$, Trace of matrix is just the diagonals
Algorithm:

- Compute gradient at each point in the image
- Compute H matrix for a window centered at every pixel in the image
$$H = \sum_{(x,y) \in W} w_{x,y} \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$
- Compute cornerness score: $\det(H) - \kappa (\text{tr}(H))^2$
- Find points whose surrounding window gave large corner response ($R >$ threshold)
- Take the points of local maxima, i.e. perform non-max suppression

Note: We cannot just find corners through those with high gradients in x & y because there are other image patches that have high gradients in both. (e.g. diagonal lines)

Non-Maximum Suppression: Iteratively search for max values, then zero out everything else
Hyperparameter: Size of window
Size of window too large → We may unintentionally remove some corners
Size of window is too small → We may include too many local maximum where in fact they are supposed to be from the same region

Adaptive Non-Maximum Suppression: Pick corners which are both local maxima and whose response is significantly greater (e.g. 10%) than all neighbouring local maxima within some radius r
Rationale: Simple search of maxima will lead to mostly interest points with high contrast (White to Black) which may not be optimal since we may want to have corners from low contrast areas as well. By doing this, we will have a better spread of corners around the image

Algorithm / Method		Invariance/Equivariance				
		Translation	Image Flip (Mirror x/y-axis)	Rotation	Scaling (proportionate)	Additive Intensity (Photometric)
Detectors	Harris Corners	✓	✓	✓		✓
	Laplacian of Gaussian	✓	✓	✓		✓
	SIFT	✓	✓	✓	✓	✓
Descriptors	MOPS	✓				
	GIST					✓
	SIFT	✓		✓	✓	✓
	Weighted Intensity	✓				✓
	Weighted Gradients	✓	✓	✓		✓

Equivariance/Invariance of Harris Corners		
Transformation	Equivariant	Invariant
Translation	True, because derivatives H is obtained through convolution which is translation equivariant	True, eigenvalues for cornerness scores is also based on derivatives
Rotation	True, eigenvectors represent the direction of max/min change in appearance, so rotate with patch	True, eigenvalues represent the corresponding magnitude of max/min change so it is constant
Photometric Transformation	Additive Changes (Changes in overall brightness) – True Scaling of Intensity (Changes in contrast) – False	
Scaling	False – Since if we scale the image up, we could detect many more points as keypoints instead	False

Laplacian of Gaussian: Highest response when the signal has the same characteristic scale as the Gaussian. (i.e. has similar shape and width of the Gaussian). It is a "blob" detector which responds strongly to circles.
Note: We can make use of difference of Gaussians to approximate this, i.e. take a bigger Gaussian to minus a smaller Gaussian as a cheaper approximation
Size of the filter: By changing the size of the filter, we will be able to detect key points of differing scale.

Descriptors		
Possible Descriptor	Benefits	Issues

Intensity	Good if geometry and appearance unchanged	If lighting changes, then we will not be able to detect these points
Gradients	Invariant to absolute intensity values	Shifting it by a little through translation and it will make all the descriptors very different
Colour Histogram	Invariant to scale (assuming the histogram is normalised) and rotation	Not sensitive to spatial layout because there could another image that has similar colour histogram but different layout of the pixels
Spatial Histogram	Retains rough spatial layout and some invariance to deformations	It is invariant to slight rotation but since it's a grid, if we rotate it too much, we no longer can find the points anymore

Histograms:
Choosing of range:

- Grids of entire domain → Fast → Only applicable when dimensionality is low
- Clustering → Slower → But can quantize data in higher dimensions

Bin Size:

- Few Bins → Less Data required → But coarser representation
- Many bins → More data required → But finer representation

Comparison between histograms: Histogram Intersection/ Euclidean could be faster but Chi-Squared often works better
Issue with Histogramming: No localisation and sensitivity to spatial layout

Multi-Scale Oriented Patches (MOPS): Local Descriptor

- Take a 40×40 window around keypoint (eg. Harris corner)
- Further split the window into 5×5 cells (subsampling → absorbs localization errors/increase tolerance)
- Rotate to horizontal (Find dominant gradient and rotate all cells to same direction) → rotation invariant
- Normalise intensity value within the window ($\frac{I - \mu}{\sigma}$) → robust to photometric variations

Wavelet transform on 8×8 patch to get a 64-dimensional descriptor (similar to PCA)

Note: Rotational Invariant, some photometric invariance but Not scale invariant
GIST Descriptors: Global Descriptor

- Divide entire image (patch) into 4×4 cells
- Apply Gabor filter (filter bank of directional edge detectors) → each pixel has N descriptors
- Compute filter response averages for each cell → each cell has N descriptors
- Size of descriptor is $4 \times 4 \times N$, where N is size of filter bank

Scale Invariant Feature Transform (SIFT): Note that the detector and descriptor can be used separately
SIFT Detector

- Multi-scale extrema detection – Scale invariance
Makes use of the LoG to detect blobs at different scales. Makes use of difference between Gaussians to estimate LoG effect
- Keypoint Localization
Refine location of keypoints to sub-pixel accuracy via Taylor series expansion across the octaves.
- Orientation assignment / rotation – Rotation invariance
Note that edges are computed through the Hessian of the DoG response image, retaining maxima only if ratio of max to min eigenvalue of the Hessian is above some threshold

Peak Threshold Parameter: For the DoG response filtering (Higher is better)

- Increase in threshold → Lesser maxima will remain

Edge Threshold Parameter: For edge response image (Lower is better)

- Decrease in threshold → Lesser edges detected → This is good because we only want corners and the ratio of max/min for corners are small whereas it is large for edges

SIFT Descriptor

- Take 16×16 window around detected keypoint from image at scale matching to keypoint. Partition window into a 4×4 grid of cells (gives some sensitivity to spatial layout)
- Compute gradient orientations for each pixel, reweight magnitudes according to a Gaussian centered on keypoint and discard pixels with low magnitude
- Of the remaining edge orientations, create a histogram with 8 orientation bins for each cell. To be rotation invariant, "shift" histogram binning by the dominant orientation. Collapse into vector ($6 \times 8 = 128$ dims)
- Normalise vector to unit length, clamp values based on threshold, re-normalize again for final descriptor

Pros	
Highly Robust	Invariant to Scale and Rotation (not strictly)
Quick and Efficient	Can handle changes in viewpoint (Up to 60 degree out of plane rotation)
	Can handle significant changes in illumination (Sometimes even day vs night)

Matching Features:

- L2 Distance:** $\|f_1 - f_2\|$ – Can give small distances for ambiguous (incorrect matches)
- Ratio Distance:** $\|f_1 - f_2\| / \|f_1 - f_3\|$ – Makes use of the ratio with the second best to ensure that we match feature pairs that are distinct

Note: We can threshold the ratio test as well to reduce the number of incorrect matches. Therefore, only those that are unique (small values) will be kept. If we increase threshold too much, number of correct matches that we get per increase will decrease because we are allowing for more ambiguous matches.

Terms	Definition
True Positive (TP)	Detected keypoint is matched correctly with its corresponding keypoint. (Each match pair is considered 1 count)
False Positive (FP)	Detected keypoint is matched incorrectly with another keypoint (Each match pair is considered 1 count)
True Negative (TN)	Keypoints which should not be matched since the corresponding keypoints were not found in the other image

False Negative (FN)	Matches that we missed. Corresponding keypoints which are detected in both images but not matched. Image (Each match pair is considered 1 count)
Precision	$TP / (TP + FP)$ – Tells us how accurate the feature pairs that are declared as matches Increase threshold → Wont know whether will increase precision
Recall	$TP / (TP + FN)$ – Tells us whether the algorithm can identify all the actual pairs of features that can be found Increase threshold → Increase recall because we will have more samples recognised as positive.
Specificity	$TN / (TN + FP)$ – Tells us whether the algorithm can correctly disregard the features which are not part of any pair Increase threshold → Decrease specificity because we will have more samples recognised as positive → More FP → Denom ↑

Solving Least Squares Conditions:

- $A^T A$ should be invertible
- $\det(A^T A)$ should not be too small, $A^{-1} = \text{adj}(\text{unc}(A)) / \det(A) \rightarrow \lambda_1$ and λ_2 should not be too small
- $A^T A$ should be well conditioned → λ_1 / λ_2 should not be too large (λ_1 is the larger value)

Automatic Scale Selection
When looking for keypoints, find the scale that give local maximum of f :
• **Harris Operator (R):** For the same keypoint, go through varying scale of window and calculate R . Use the scale that maximises R value.
• **Laplacian of Gaussians/Blob Detector:** Similar idea. Alternatively, instead of scaling filter size, we can scale the image and fix the filter size but also iterate through the various sigma value.

Homography
Applying a Homography:

- Convert to homogeneous coordinates. $p = [x \ y \ 1]^T \Rightarrow P = [x \ y \ 1]$
- Multiply by the homography matrix $P' = H \cdot P$
- Convert back to heterogeneous coordinates $P' = [x' \ y' \ w']^T \Rightarrow p' = [\frac{x'}{w'} \ \frac{y'}{w'}]$

Solving for the H matrix: Note that we need 4 pairs of points (8 points total) to solve the homography and we only have 8 degrees of freedom because we can divide by h_3 throughout and it gets absorbed.

$$\begin{aligned} A_1 h &= 0 \\ h_2 x x' + h_3 x y' + h_4 x x' - h_1 x - h_2 y - h_3 &= 0 \dots (1) \\ h_4 x y' + h_5 y y' + h_6 x y' - h_1 x - h_2 y - h_3 &= 0 \dots (2) \\ A_i &= \begin{bmatrix} -x & -y & -1 & 0 & 0 & 0 & x x' & x y' & x' \\ 0 & 0 & 0 & -x & -y & -1 & x y' & y y' & y' \end{bmatrix} \\ h &= [h_1 \ h_2 \ h_3 \ h_4 \ h_5 \ h_6 \ h_7 \ h_8 \ h_9] \end{aligned}$$

Solve for H using DLT:
[Given p_i, p'_i] solve for H such that $P'_i = H \cdot P_i \ \forall i$

- For each correspondence, create 2×9 matrix A_i
- Concatenate into single $2n \times 9$ matrix A
- Compute SVD of $A = U \Sigma V^T$
- Store singular vector of smallest singular value $h = v_1$
- Reshape h to get H

Cons of using DLT: Note that the points here works with applicability of Homography

- Only applicable when the actual transformation between the two images is linear. If that is the case then this may not work out well
- Sensitive to scaling in pixel space. (We can add a normalization step)
- Does not deal well with outliers (i.e. bad matches). Can be more robust with RANSAC

Applicability of Homography:

- The scene is **planar** or **approximately planar** (i.e. scene is very far away or has small (relative) depth variation.
- Scene is captured under camera rotation only (no translation or pose change)

Normalised DLT:

- Normalisation of p :** Compute a similarity transformation T , consisting of a translation and scaling, that takes points p_i to a new set of points \tilde{x}_i such that the centroid of the points \tilde{x}_i is (0, 0) and their average distance from the origin is $\sqrt{2}$
- Normalisation of p' :** Compute a similar transformation T' for the points in the second image, transforming points p'_i to \tilde{p}'_i
- DLT:** Apply standard DLT to the correspondences $\tilde{p}_i \leftrightarrow \tilde{p}'_i$ to obtain a homography \tilde{H}
- Denormalization:** Set $H = T'^{-1} \tilde{H} T$

RANSAC Algorithm:
Loop until total number of iterations reached

- Randomly get (min 4) correspondences
- Compute H using DLT
- Count inliers. Note that we may have to consider squaring the distance to ensure that we are getting inliers from both sides.
- Keep H if largest numbers of inliers

For the best H with the most inliers, recompute H again using all the new inliers.

Parameters:

- δ – Distance threshold to consider as an inlier
Should be large enough so that noisy inlier points will still lie within the threshold
- N – Number of iterations to run RANSAC loop
$$N = \frac{\log(1 - p)}{\log(1 - (1 - \epsilon)^2)}$$

ϵ – Prob that point is an outlier, s – Number of points we draw per iteration for fitting parameters, N – Number of iterations in loop, p – probability that we want to ensure that at least one set of points which we sampled does not contain any outliers

Optical Flow

Flow Methods	Lucas-Kanade Flow	Horn-Shunck Flow
	Can be used when image motion is large (Pyramid approach)	Suitable for video when image motion is small
Assumptions	Brightness Constancy – Colour Constancy and Small Motion $I_x u + I_y v + I_t = 0$ Make use of $\partial x = [-1 \ 0 \ 1], \partial y = [-1 \ 0 \ 1]$	
Method	Constant Flow – Assume surrounding pixels (e.g. 5×5) have the same displacement Featured based approach. Will only track visual features (corners, textured areas)	Smooth Flow Field – Neighbouring pixels should have similar flow to the current pixel (don't have to be same) Recovers image motions at every pixel based on spatio-temporal image brightness variation
Type of Flow Field	Sparse motion fields due to the flow fields only detected at visual features but it is more robust at tracking	Dense motion fields due to flow fields detected at every pixel but it is sensitive to appearance variations
Issues	Aperture Problem – If we only have a small visible patch and a line, we wont be able to know what is the exact motion. Aliasing – If there are many pixels with the same intensity, when the motion is too big, there could be wrong "correspondence". Solution: Coarse-to-Fine Estimation. Edges – Edges that are detected will have perpendicular flows	λ small → Care more about smoothness λ large → Care more about brightness Swirl Patterns – Due to strong smoothing constraints "Apparent Motions" – Could be due to shadows

Tracking
Lucas Kanade (Additive) Alignment:

- Warp image, $I(W(x; p))$
- Compute error image $|T(x) - I(W(x; p))|$
- Compute gradient of the warped image $\nabla I(x')$
- Evaluate Jacobian $\partial W / \partial p$
- Compute Hessian approximation $H = \sum_i [\nabla I \partial W / \partial p]_i^T [\nabla I \partial W / \partial p]_i$
- Compute $\Delta p = H^{-1} \sum_i [\nabla I \partial W / \partial p]_i^T [T(x) - I(W(x; p))]$
- Update parameters $p \leftarrow p + \Delta p$

KL T Algorithm for Feature Tracking:

- Find corners satisfying $\min(\lambda_1, \lambda_2) > \lambda$ for the Hessian
- Loop over corners

Compute displacement to next frame using the Lucas-Kanade alignment method. Store displacement of each corner, update corner position. (optional) Add more corner points every M frames using step 1

- Returns long trajectories for each corner point

Tracking Parameters:
 η – Controls how much we want the new value to affect our template. When η is large, it means that we want more of the new values.

Tracking Algorithms differ by:		
Property	Approaches	
Representing Candidates	Naïve: Grey scale/ RGB Better: Gradient Features, Histogramming, Deep Features	
Search for Candidates	Naïve: Exhaustive set of locations. (Computation and efficiency concerns) Better: Limit search space based on previous results.	
Solving for mode-finding problem	Naïve: Take a simple global max (But could be hard to distinguish sometimes) Better: Leverage previous locations to help with local max	
Updating target's template	Naïve: Keep fixed, don't update but there could be appearance changes Better: Update as the track progresses and using η as a tuning parameter	

Issues with Tracking:

- Initialisation** – Often done manually. We can make use of background subtraction, detection can also be used
- Catastrophic errors** – Occlusions & clutter, exit of frame, multiple objects
- Drifting** – Due to accumulation of small errors over time

MOSSE Filter for Tracking:

- Initialise on frame I_0 with marked ROI x_0 (DFT → Optimise kernel g)
- Apply kernel g to frame I_k (DFT)
- Find the local maxima to locate target bounding box x_k (IDFT)
- Update kernel with new bounding box x_k (DFT → Optimise kernel g)

Note: Initialisation is done on several perturbed version of initial ROI to add robustness to first kernel g . Kernel is solved in Fourier domain via discrete Fourier transform (DFT). Localisation is in the image domain, so convert cross-correlation results back with inverse DFT (IDFT).

Evaluation Metric for Tracking:

- Accuracy** – Talks about how well the tracker's bounding box overlap with ground truth bounding box
Can make use of $|B_{gt} \cap B_p| / |B_{gt} \cup B_p|$ as one of the metric. We can average it over frames in sequence for per-sequence accuracy. $[0, 1]$ where 1 is the best
- Robustness** – Talks about how many times the tracker fails (i.e. lose the target) (0 is ideal)
We can re-initialise the tracker (manually) each time it has failed to resume tracking. We will average over multiple tracking runs on a sequence to see how many times it fails on average.